

Homework 2

Shao-Ting Chiu (UIN:433002162)

10/12/22

Table of contents

Homework Description	1
Computational Enviromnent Setup	2
Third-party libraries	2
Version	2
Problem 3.6 (Python Assignment)	3
Problem 4.2	5
(a)	6
(b)	6
(c)	8
Problem 4.3	8
(a)	9
(b)	11
Problem 4.4	12
(a)	12
(b)	14
(c)	16
Problem 4.8 (Python Assignment)	18
(a)	18
(b)	19
(c)	20
(d)	23
References	24

Homework Description

- Course: ECEN649, Fall2022
- Problems from the book:
- 3.6 (10 pt)
 - 4.2 (10 pt)
 - 4.3 (10 pt)

4.4 (10 pt)

4.8 (20 pt)

- Deadline: Oct. 12th, 11:59 am

Computational Environment Setup

Third-party libraries

```
1 %matplotlib inline
2 import sys # system information
3 import matplotlib # plotting
4 import scipy.stats as st # scientific computing
5 import pandas as pd # data managing
6 import numpy as np # numerical computation
7 import numba
8 import sklearn as sk
9 from numpy import linalg as LA
10 import scipy as sp
11 import scipy.optimize as opt
12 import sympy as sp
13 import matplotlib.pyplot as plt
14 from numpy.linalg import inv, det
15 from numpy.random import multivariate_normal as mvn
16 from numpy.random import binomial as binom
17 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA #problem 4.8
18 from sklearn.model_selection import train_test_split
19 # Matplotlib setting
20 plt.rcParams['text.usetex'] = True
21 matplotlib.rcParams['figure.dpi']= 300
22 np.random.seed(20221011)
```

Version

```
1 print(sys.version)
2 print(matplotlib.__version__)
3 print(sp.__version__)
4 print(np.__version__)
5 print(pd.__version__)
6 print(sk.__version__)
```

3.8.14 (default, Sep 6 2022, 23:26:50)
[Clang 13.1.6 (clang-1316.0.21.2.5)]
3.3.1

1.6.2
 1.19.1
 1.1.1
 1.1.2

Problem 3.6 (Python Assignment)

Using the synthetic data model in Section A8.1 for the homoskedastic case with $\mu_0 = (0, \dots, 0)$, $\mu_1 = (1, \dots, 1)$, $P(Y = 0) = P(Y = 1)$, and $k = d$ (independent features), generate a large number (e.g., $M = 1000$) of training data sets for each sample size $n = 20$ to $n = 100$, in steps of 10, with $d = 2, 5, 8$, and $\sigma = 1$. Obtain an approximation of the expected classification error $E[\epsilon_n]$ of the nearest centroid classifier in each case by averaging ϵ_n , computed using the exact formula (3.13), over the M synthetic training data sets. Plot $E[\epsilon_n]$ as a function of the sample size, for $d = 2, 5, 8$ (join the individual points with lines to obtain a smooth curve). Explain what you see.

- The formula in Braga-Neto (2020, 56, Eq. 3.13)
- $\epsilon_n = \frac{1}{2} \left(\Phi \left(\frac{a_n^T \hat{\mu}_0 + b_n}{\|a_n\|} \right) + \Phi \left(-\frac{a_n^T \hat{\mu}_1 + b_n}{\|a_n\|} \right) \right)$
 - $\mu_0 = (0, \dots, 0)$; $\hat{\mu}_0 = \frac{1}{N_0} \sum_{i=1}^n X_i I_{Y_i=0}$
 - $\mu_1 = (1, \dots, 1)$; $\hat{\mu}_1 = \frac{1}{N_1} \sum_{i=1}^n X_i I_{Y_i=1}$
 - $a_n = \hat{\mu}_1 - \hat{\mu}_0$
 - $b_n = -\frac{(\hat{\mu}_1 - \hat{\mu}_0)^T (\hat{\mu}_1 + \hat{\mu}_0)}{2} \mathbf{1}$
- As shown in Figure 1, the error rate converges to optimal error as the sample size increases.

```

1 def norml(v):
2     return LA.norm(v, 2)
3
4 def get_an(hm0, hm1):
5     return hm1 - hm0
6
7 def get_bn(hm0, hm1):
8     return -float((hm1 - hm0).T @ (hm1 + hm0))/2
9
10 def epsilon(hmu0, hmu1, mu0, mu1, p0=0.5):
11     p1 = 1-p0

```

¹ All, in Example 3.4 there is a negative sign missing. The value of bn is $-(m1 - m0)^T(m1 + m0)/2$. — [Ulisses on Slack](#)

```

12     an = get_an(hmu0, hmu1)
13     bn = get_bn(hmu0, hmu1)
14     epsilon0 = st.norm.cdf( (float(an.T* mu0) + bn)/norml(an))
15     epsilon1 = st.norm.cdf(- (float(an.T*mu1)+ bn)/norml(an))
16     return p0*epsilon0 + p1*epsilon1
17
18 class GaussianDataGen:
19     def __init__(self, n, d, s=1, mu=0):
20         self.n = n
21         self.d = d
22         self.mu = np.matrix(np.ones(d) * mu).T
23         self.s = s
24         self.cov = self.get_cov()
25
26     def get_cov(self):
27         return np.identity(self.d) * self.s
28
29     def sample(self):
30         data = np.random.normal(self.mu[0][0], self.s, size= (self.d, self.n))
31         hmuV = np.mean(data, axis=1)
32         return np.matrix(hmuV).T
33
34 def cal_eps(dg0, dg1, p0=0.5):
35     hmuV0 = dg0.sample()
36     hmuV1 = dg1.sample()
37     mu0 = np.matrix(np.zeros(dg0.d)).T
38     mu1 = np.matrix(np.ones(dg1.d)).T
39     return epsilon(hmuV0, hmuV1, mu0, mu1,p0=0.5)
40 cal_eps_func = np.vectorize(cal_eps)
41
42 def exp_try_nd(n, d, s=1,M=1000):
43     gX0 = GaussianDataGen(n=n, d=d, s= s,mu=0)
44     gX1 = GaussianDataGen(n=n, d=d, s= s, mu=1)
45     #eps = cal_eps_func([gX0]*M, gX1)
46     eps = [cal_eps_func(gX0, gX1) for i in range(0,M)]
47     return np.mean(eps)
48 exp_try_nd_func = np.vectorize(exp_try_nd)
49
50 def bayes_ncc(mu0, mu1):
51     return st.norm.cdf(- norml(mu1-mu0)/2)
52
53 M = 1000
54 ns = np.arange(20,100, 10)
55 s = 1

```

```

56 dres = {2:[],5:[],8:[]}
57
58 for k in dres.keys():
59     dres[k]= exp_try_nd_func(ns,k)
60
61 # Optimal error
62 opts = [bayes_ncc(np.zeros(k), np.ones(k)) for k in dres.keys()]
63
64 fig, ax = plt.subplots()
65 sts = ["-", "--", ":"]
66 for (i, k) in enumerate(dres.keys()):
67     ax.plot(ns, dres[k], 'o',label="d={}".format(k))
68     ax.axhline(y= opts[i], label="Optimal d = {}".format(k), color='k', linestyle=sts[i])
69
70 ax.set_xlabel("n")
71 ax.set_ylabel("$E[\epsilon_n]$")
72 ax.legend();

```




Figure 1: Error of E_n

Problem 4.2

A common method to extend binary classification rules to K classes, $K > 2$, is the *one-vs-one approach*, in which $\frac{K(K-1)}{2}$ classifiers are

trained between all pairs of classes², and a majority vote of assigned labels is taken.

(a)

Formulate a multiclass version of parametric plug-in classification using the one-vs-one approach.

Let $\psi_{i,j}^*$ be a one-one classifiers that $i \neq j$, and $\{(i, j) | i \in [1, k], j \in [1, k], i \neq j\}$. For K classes, there are $K(K-1)$ classifiers; for each classifier $\psi_{i,j}^*$ and $x \in R^d$,

$$\psi_{ij,n}^* = \begin{cases} 1, & D_{ij,n}(x) > k_{ij,n} \\ 0, & \text{otherwise} \end{cases} \quad (1)$$

where

- $D_{ij,n}(x) = \ln \frac{p(x|\theta_{i,n})}{p(x|\theta_{j,n})}$
- $k_{ij,n} = \ln \frac{P(Y=j)}{P(Y=i)}$
- Noted that feature-label distribution is expressed via a family of PDF $\{p(x|\theta_i) | \theta \in \Theta \subseteq R^m\}$, for $i = 1, \dots, K$.
- $\psi_{ij,n}^* = 1 \otimes \psi_{ji,n}^*$. These two are similar classifiers with inverted outcome.

Let $\psi_{i,n}^* = \sum_{j \neq i} \psi_{ij,n}^*$, and the one-vs-one classifier is

$$\psi_n^*(x) = \arg \max_{k=1, \dots, K} \psi_{k,n}^* \quad (2)$$

(b)

Show that if the threshold $k_{ij,n}$ between classes i and j is given by $\ln \frac{\hat{c}_j}{\hat{c}_i}$, then the one-vs-one parametric classification rule is equivalent to the simple decision.

$$\psi_n(x) = \arg \max_{k=1, \dots, K} \hat{c}_k p(x|\theta_{k,n}), x \in R^d$$

(For simplicity, you may ignore the possibility of ties.)

$$\ln \frac{p(x|\theta_{i,n})}{p(x|\theta_{j,n})} > k_{ij,n} = \ln \frac{\hat{c}_j}{\hat{c}_i} \quad (3)$$

$$\ln p(x|\theta_{i,n}) - \ln p(x|\theta_{j,n}) > \ln \hat{c}_j - \ln \hat{c}_i \quad (4)$$

$$\hat{c}_i p(x|\theta_{i,n}) > \hat{c}_j p(x|\theta_{j,n}) \quad (5)$$

²

Also, in Problem 4.2, the number of classifiers is $K(K-1)/2$ not $K(K-1)$ — [Ulisses on TAMU Slack](#)

$$\psi_{ij,n}^* = \begin{cases} 1, & \hat{c}_i p(x|\theta_{i,n}) > \hat{c}_j p(x|\theta_{j,n}) \\ 0, & \text{otherwise} \end{cases} \quad (6)$$

$$= I_{\hat{c}_i p(x|\theta_{i,n}) > \hat{c}_j p(x|\theta_{j,n})} \quad (7)$$

Then,

$$\psi_{i,n}^* = \sum_{j \neq i} \psi_{ij,n}^* \quad (8)$$

$$= \sum_{j \neq i} I_{\hat{c}_i p(x|\theta_{i,n}) > \hat{c}_j p(x|\theta_{j,n})} \quad (9)$$

$$\psi_n^*(x) = \arg \max_{k=1,\dots,K} \psi_{k,n}^* \quad (10)$$

$$= \arg \max_{k=1,\dots,K} \sum_{j \neq i} \psi_{kj,n}^* \quad (11)$$

$$= \arg \max_{k=1,\dots,K} \sum_{j \neq i} I_{\hat{c}_k p(x|\theta_{k,n}) > \hat{c}_j p(x|\theta_{j,n})} \quad (12)$$

$$(13)$$

Let $\psi_n^*(x) = \kappa$, that means $\psi_{\kappa,n}^*$ is the maximum among $\{\psi_{j,n}^* | j = (1, \dots, K)\}$. Assume there is an $s \neq \kappa$ s.t. $\hat{c}_s p(x|\theta_{s,n}) > \hat{c}_\kappa p(x|\theta_{\kappa,n}) > \text{the rests}$. Thus, $I_{\hat{c}_s p(x|\theta_{s,n}) > \hat{c}_\kappa p(x|\theta_{\kappa,n})} = 1$

$$\psi_{s,n}^* = \sum_{j \neq s} I_{\hat{c}_s p(x|\theta_{s,n}) > \hat{c}_j p(x|\theta_{j,n})} = \underbrace{\sum_{j \neq s; j \neq \kappa} I_{\hat{c}_s p(x|\theta_{s,n}) > \hat{c}_j p(x|\theta_{j,n})}}_{=a} + \underbrace{I_{\hat{c}_s p(x|\theta_{s,n}) > \hat{c}_\kappa p(x|\theta_{\kappa,n})}}_{=1}$$

$$\psi_{\kappa,n}^* = \sum_{j \neq \kappa} I_{\hat{c}_\kappa p(x|\theta_{\kappa,n}) > \hat{c}_j p(x|\theta_{j,n})} = \underbrace{\sum_{j \neq \kappa; j \neq s} I_{\hat{c}_\kappa p(x|\theta_{\kappa,n}) > \hat{c}_j p(x|\theta_{j,n})}}_{=a} + \underbrace{I_{\hat{c}_\kappa p(x|\theta_{\kappa,n}) > \hat{c}_s p(x|\theta_{s,n})}}_{=0}$$

where a is a nonnegative number. That means

$$\psi_{\kappa,n}^* < \psi_{s,n}^* \quad (14)$$

$\psi_{\kappa,n}^*$ is not the maximum. Equation 14 is contradict to the statement that $\psi_n^*(x) = \kappa$. In conclusion, $\hat{c}_k p(x|\theta_{\kappa,n})$ is the maximum if $\psi_n(x) = k$.

(c)

Applying the approach in items (a) and (b), formulate a multiclass version of Gaussian discriminant analysis. In the case of multiclass NMC, with all thresholds equal to zero, how does the decision boundary look like?

For Gaussian discriminant analysis, the discriminant is defined as

$$\hat{D}_{ij}^*(x) = \frac{1}{2}(x - \hat{\mu}_i)^T \hat{\Sigma}_i^{-1}(x - \hat{\mu}_i) - \frac{1}{2}(x - \hat{\mu}_j)^T \hat{\Sigma}_j^{-1}(x - \hat{\mu}_j) + \frac{1}{2} \ln \frac{\det(\hat{\Sigma}_i)}{\det(\hat{\Sigma}_j)}$$

$$\psi_{ij,n}^*(x) = I_{\hat{D}_{ij}^*(x) > 0} \quad (15)$$

$$\psi_n^* = \arg \max_{k=1, \dots, K} \hat{c}_k p(x | \hat{\mu}_{k,n}, \hat{\Sigma}_{k,n}) = \arg \max_{k=1, \dots, K} \text{Normal}(x; \hat{\mu}_{k,n}, \hat{\Sigma}_{k,n})$$

For NMC case, suppose there are 3 classes with $d = 2$.

- $K = 3$
- Number of classifiers: $\frac{3 \cdot 2}{2} = 3$
 - $\psi_{1,2}^*, \psi_{1,3}^*, \psi_{2,3}^*$

$$\text{boundary} = \begin{cases} \hat{\Sigma}^{-1}(\hat{\mu}_1 - \hat{\mu}_2) \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + (\hat{\mu}_2 - \hat{\mu}_1)^T \hat{\Sigma}^{-1} \left(\frac{\hat{\mu}_1 + \hat{\mu}_2}{2} \right) = 0 \\ \hat{\Sigma}^{-1}(\hat{\mu}_1 - \hat{\mu}_3) \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + (\hat{\mu}_3 - \hat{\mu}_1)^T \hat{\Sigma}^{-1} \left(\frac{\hat{\mu}_1 + \hat{\mu}_3}{2} \right) = 0 \\ \hat{\Sigma}^{-1}(\hat{\mu}_2 - \hat{\mu}_3) \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + (\hat{\mu}_3 - \hat{\mu}_2)^T \hat{\Sigma}^{-1} \left(\frac{\hat{\mu}_2 + \hat{\mu}_3}{2} \right) = 0 \end{cases} \quad (16)$$

Problem 4.3

Under the general Gaussian model $p(x|Y=0) \sim \mathcal{N}_d(\mu_0, \Sigma_0)$ and $p(x|Y=1) \sim \mathcal{N}_d(\mu_1, \Sigma_1)$, the classification error $\epsilon_n = P(\psi_n(X) \neq Y | S_n)$ of *any* linear classifier in the form

$$\psi_n(x) = \begin{cases} 1, & a_n^T x + b_n > 0, \\ 0, & \text{otherwise} \end{cases} \quad (17)$$

(examples discussed so far include LDA and its variants, and the logistic classifier) can be readily computed in terms of Φ (the CDF




Figure 2: Homoskedastic cases

of a standard normal random variable), the classifier parameters a_n and b_n , and the distributional parameters $c = P(Y = 1)$, μ_0 , μ_1 , Σ_0 , and Σ_1 .

(a)

Show that

$$\epsilon_n = (1 - c) \Phi \left(\frac{a_n^T \mu_0 + b_n}{\sqrt{a_n^T \Sigma_0 a_n}} \right) + c \Phi \left(-\frac{a_n^T \mu_1 + b_n}{\sqrt{a_n^T \Sigma_1 a_n}} \right)$$

Hint: the discriminant $a_n^T x + b_n$ has a simple Gaussian distribution in each class.

From Braga-Neto (2020, Eq. 2.34),

$$\epsilon^* = \underbrace{P(Y = 0)}_{=1-c} \epsilon^0[\psi^*] + \underbrace{P(Y = 1)}_{=c} \epsilon^1[\psi^*]$$

$$\epsilon^0[\psi^*] = P(a_n^T x + b_n > 0 | Y = 0) \quad (18)$$

$$(19)$$

Use the affine property of Gaussian distribution described in Braga-Neto (2020) [pp. 307. G4]³.

$$\bullet \quad a_n^T x + b_n | Y = 0 \sim N(a_n^T \mu_0 + b_n, \underbrace{a_n^T \Sigma_0 a_n}_{\sigma^2})$$

$$\epsilon^0[\psi^*] = 1 - P(a_n^T x + b_n \leq 0 | Y = 0) \quad (20)$$

$$= 1 - \Phi \left(\frac{0 - (a_n^T \mu_0 + b_n)}{\sqrt{a_n^T \Sigma_0 a_n}} \right) \quad (21)$$

$$= 1 - \Phi \left(-\frac{a_n^T \mu_0 + b_n}{\sqrt{a_n^T \Sigma_0 a_n}} \right) \quad (22)$$

$$= \Phi \left(\frac{a_n^T \mu_0 + b_n}{\sqrt{a_n^T \Sigma_0 a_n}} \right) \quad (23)$$

Similarly,

$$\epsilon^1(\psi^*) = P(a_n^T x + b_n < 0 | Y = 1) \quad (24)$$

$$= \Phi \left(\frac{0 - (a_n^T \mu_1 + b_n)}{\sqrt{a_n^T \Sigma_1 a_n}} \right) \quad (25)$$

$$= \Phi \left(-\frac{a_n^T \mu_1 + b_n}{\sqrt{a_n^T \Sigma_1 a_n}} \right) \quad (26)$$

Combining together,

$$\epsilon^* = \underbrace{P(Y = 0)}_{=1-c} \epsilon^0[\psi^*] + \underbrace{P(Y = 1)}_{=c} \epsilon^1[\psi^*] \quad (27)$$

$$= (1 - c) \epsilon^0[\psi^*] + c \epsilon^1[\psi^*] \quad (28)$$

$$= (1 - c) \Phi \left(\frac{a_n^T \mu_0 + b_n}{\sqrt{a_n^T \Sigma_0 a_n}} \right) + c \Phi \left(-\frac{a_n^T \mu_1 + b_n}{\sqrt{a_n^T \Sigma_1 a_n}} \right) \quad (29)$$

³Any affine transformation $f(x) = AX + B$ of a Gaussian is a Gaussian. That is, A is a nonsingular matrix, and B is a vector. If $X \sim N(\mu, \Sigma)$, $a^T X + b \sim N(a^T \mu + b, a^T \Sigma a)$. (ardianumam 2017)

(b)

Compute the errors of the NMC, LDA, and DLDA classifiers in Example 4.2 if $c = 1/2$,

$$\mu_0 = \begin{bmatrix} 2 \\ 3 \end{bmatrix}, \mu_1 = \begin{bmatrix} 6 \\ 5 \end{bmatrix}, \Sigma_0 = \begin{bmatrix} 1 & 1 \\ 1 & 2 \end{bmatrix}, \text{ and } \Sigma_1 = \begin{bmatrix} 4 & 0 \\ 0 & 1 \end{bmatrix}$$

Which classifier does the best?

As shown in Table 1, NMC has the lowest Bayes error.

```
1 def epsilon_general(an, bn, c, mu0, mu1, sig0, sig1):
2
3     e0 = st.norm.cdf(\
4         float(an.T @ mu0 + bn)/\
5         np.sqrt(float(an.T @ sig0 @ an)))
6     e1 = st.norm.cdf(\
7         -float(an.T @ mu1 + bn)/\
8         np.sqrt(float(an.T @ sig1 @ an)) )
9
10    return (1-c)*e0 + c*e1
11
12    truth = {
13        "c": 0.5,
14        "mu0": np.matrix([[2],[3]]),
15        "mu1": np.matrix([[6],[5]]),
16        "sig0": np.matrix([[1,1],[1,2]]),
17        "sig1": np.matrix([[4,0],[0,1]]),
18    }
19
20    meth = {
21        "NMC":{
22            "an": np.matrix([[4],[2]]),
23            "bn": -24
24        },
25        "LDA":{
26            "an": 3/7*np.matrix([[5],[3]]),
27            "bn": -96/7
28        },
29        "DLDA":{
30            "an": 2/5*np.matrix([[6],[5]]),
31            "bn": -88/5
32        }
33    }
34
35    berrors = np.zeros(len(meth.keys()))
```

Table 1: Bayes Errors of NMC, LDA and DLDA

	Method	Bayes Error
0	NMC	0.084775
1	LDA	0.085298
2	DLDA	0.087606

```

36
37 for (i,k) in enumerate(meth.keys()):
38     berrors[i] = epsilon_general(**meth[k], **truth)
39
40 pd.DataFrame({"Method": list(meth.keys()),\
41              "Bayes Error": berrors}).sort_values(\
42              ["Bayes Error"], ascending=[1])

```

Problem 4.4

Even in the Gaussian case, the classification error of quadratic classifiers in general require numerical integration for its computation. In some special simple cases, however, it is possible to obtain exact solutions. Assume a two-dimensional Gaussian problem with $P(Y = 1) = \frac{1}{2}$, $\mu_0 = \mu_1 = 0$, $\Sigma_0 = \sigma_0^2 I_2$, and $\Sigma_1 = \sigma_1^2 I_2$. For definiteness, assume that $\sigma_0 < \sigma_1$.

(a)

Show that the Bayes classifier is given by

$$\psi^*(x) = \begin{cases} 1, & \|x\| > r^*, \\ 0, & \text{otherwise,} \end{cases} \quad \text{where } r^* = \sqrt{2 \left(\frac{1}{\sigma_0^2} - \frac{1}{\sigma_1^2} \right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}} \quad (30)$$

In particular, the optimal decision boundary is a circle of radius r^* .

The inverted Σ_1 and Σ_2 are⁴

4

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix}^{-1} = \frac{1}{ad - bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix} \quad (31)$$

$$\Sigma_0 = \sigma_0^2 I_2 = \begin{bmatrix} \sigma_0^2 & 0 \\ 0 & \sigma_0^2 \end{bmatrix} \quad (32)$$

$$\Sigma_0^{-1} = \frac{1}{\sigma_0^4} \begin{bmatrix} \sigma_0^2 & 0 \\ 0 & \sigma_0^2 \end{bmatrix} = \sigma_0^{-2} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \sigma_0^{-2} I_2 \quad (33)$$

$$\Sigma_1^{-1} = \sigma_1^{-2} I_2 \quad (34)$$

Use the derivation in Braga-Neto (2020, 74),

$$A_n = \begin{bmatrix} a_{11} & a_{12} \\ a_{12} & a_{22} \end{bmatrix} = \frac{-1}{2} \Sigma_1^{-1} - \Sigma_0^{-1} = \frac{-1}{2} (\sigma_1^{-2} - \sigma_0^{-2}) \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \quad (35)$$

$$b_n = \begin{bmatrix} b_{n,1} \\ b_{n,2} \end{bmatrix} = \Sigma_1^{-1} \underbrace{\mu_1}_{=0} - \Sigma_0^{-1} \underbrace{\mu_0}_{=0} \quad (36)$$

$$= \begin{bmatrix} 0 \\ 0 \end{bmatrix} \quad (37)$$

$$c = -\frac{1}{2} \ln \frac{|\Sigma_1|}{|\Sigma_0|} = \frac{-1}{2} \ln \frac{\sigma_1^4}{\sigma_0^4} = -\ln \frac{\sigma_1^2}{\sigma_0^2}$$

According to Braga-Neto (2020, Eq. 4.26), the 2-dimensional QDA decision boundary is

$$D(x) = a_{11}x_1^2 + 2a_{12}x_1x_2 + a_{22}x_2^2 + b_1x_1 + b_2x_2 + c = 0 \quad (38)$$

$$a_{11}(x_1^2 + x_2^2) = \ln \frac{\sigma_1^2}{\sigma_0^2} \quad (39)$$

$$x_1^2 + x_2^2 = 2\left(\frac{1}{\sigma_0^2} - \frac{1}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2} \quad (40)$$

$$r^* = \sqrt{x_1^2 + x_2^2} = \sqrt{2\left(\frac{1}{\sigma_0^2} - \frac{1}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}} \quad (41)$$

Noted that $\left(\frac{1}{\sigma_0^2} - \frac{1}{\sigma_1^2}\right) > 0$ because $\sigma_0 < \sigma_1$

For any point $\|x_j\| > r^*$, the discriminant (D) is larger than 0, and $\psi^*(x_j) = 1$.

(b)

Show that the corresponding Bayes error is given by

$$\epsilon^* = \frac{1}{2} - \frac{1}{2} \left(\frac{\sigma_1^2}{\sigma_0^2} - 1 \right) e^{-(1 - \frac{\sigma_0^2}{\sigma_1^2})^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}}$$

In particular, the Bayes error is a function only of the ratio of variances $\frac{\sigma_1^2}{\sigma_0^2}$, and $\epsilon^* \rightarrow 0$ as $\frac{\sigma_1^2}{\sigma_0^2} \rightarrow \infty$.

Hint: use polar coordinates to solve the required integrals analytically.

Part I: Definition of errors

$$\epsilon^0[\psi^*] = P(D^*(X) > k^* | Y = 0) \quad (42)$$

$$= P(\|x\| > r^* | Y = 0) \quad (43)$$

$$\epsilon^1[\psi^*] = P(D^*(X) \leq k^* | Y = 1) \quad (44)$$

$$= P(\|x\| \leq r^* | Y = 1) \quad (45)$$

Part II: PDF of 2D Gaussian

$$p(x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}) = \frac{1}{\sqrt{(2\pi)^2 \sigma^4}} \exp\left(-\frac{1}{2} x^T \Sigma^{-1} x\right) \quad (46)$$

$$= \frac{1}{\sqrt{(2\pi)^2 \sigma^4}} \exp\left(-\frac{1}{2} \frac{x_1^2 + x_2^2}{\sigma^2}\right) \quad (47)$$

$$= \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x_1^2 + x_2^2}{2\sigma^2}\right) \quad (48)$$

$$(49)$$

Use the polar coordination, $x_1 = r \cos \theta$ and $x_2 = r \sin \theta$. $x_1^2 + x_2^2 = r^2$. We can transform 2D gaussian into polar coordination:

$$p(r, \theta) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{r^2}{2\sigma^2}\right)$$

Part III: Integration

$$\epsilon^0[\psi^*] = \int_{\theta=0}^{\theta=2\pi} \int_{r=r^*}^{\infty} \frac{1}{2\pi\sigma_0^2} \exp\left(-\frac{r^2}{2\sigma_0^2}\right) r dr d\theta \quad (50)$$

$$= \frac{1}{2\pi\sigma_0^2} \int_{\theta=0}^{\theta=2\pi} \int_{r=r^*}^{\infty} \exp\left(-\frac{r^2}{2\sigma_0^2}\right) r dr d\theta \quad (51)$$

$$= \frac{1}{2\pi\sigma_0^2} \int_{\theta=0}^{\theta=2\pi} \sigma_0^2 \exp\left(-\frac{r_*^2}{2\sigma_0^2}\right) d\theta \quad (52)$$

$$= \exp\left(-\frac{r_*^2}{2\sigma_0^2}\right) \quad (53)$$

$$= \exp\left(-\frac{1}{\sigma_0^2(\sigma_0^{-2} - \sigma_1^{-2})} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (54)$$

$$= \exp\left(\frac{-1}{(1 - \frac{\sigma_0^2}{\sigma_1^2})} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (55)$$

$$= \exp\left(-\left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (56)$$

$$\epsilon^1[\psi^*] = \int_{\theta=0}^{\theta=2\pi} \int_{r=0}^{r=r^*} \frac{1}{2\pi\sigma_1^2} \exp\left(-\frac{r^2}{2\sigma_1^2}\right) r dr d\theta \quad (57)$$

$$= 1 - \exp\left(-\frac{r_*^2}{2\sigma_1^2}\right) \quad (58)$$

$$= 1 - \exp\left(-\frac{1}{\sigma_1^2(\sigma_0^{-2} - \sigma_1^{-2})} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (59)$$

$$= 1 - \exp\left(-\frac{1}{(\frac{\sigma_1^2}{\sigma_0^2} - 1)} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (60)$$

$$= 1 - \exp\left(-\frac{\frac{\sigma_0^2}{\sigma_1^2}}{(1 - \frac{\sigma_0^2}{\sigma_1^2})} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (61)$$

$$= 1 - \exp\left(-\frac{\sigma_0^2}{\sigma_1^2} \left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (62)$$

Part IV: Combining together

$$\epsilon^* = P(Y = 0)\epsilon^0[\psi^*] + P(Y = 1)\epsilon^1[\psi^*] \quad (63)$$

$$= \frac{1}{2}\epsilon^0 + \frac{1}{2}\epsilon^1 \quad (64)$$

$$= \frac{1}{2} \exp\left(-\left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) + \frac{1}{2} - \frac{1}{2} \exp\left(-\frac{\sigma_0^2}{\sigma_1^2} \left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (65)$$

$$= \frac{1}{2} \exp\left(-\left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) + \frac{1}{2} \quad (66)$$

$$- \frac{1}{2} \exp\left(-\left(\frac{\sigma_0^2}{\sigma_1^2} - 1\right) \left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \exp\left(-\left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (67)$$

$$= \frac{1}{2} + \frac{1}{2} \left[1 - \exp\left(-\underbrace{\left(\frac{\sigma_0^2}{\sigma_1^2} - 1\right) \left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1}}_{=-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \right] \exp\left(-\left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (68)$$

$$= \frac{1}{2} + \frac{1}{2} \left[1 - \exp\left(\ln \frac{\sigma_1^2}{\sigma_0^2}\right) \right] \exp\left(-\left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (69)$$

$$= \frac{1}{2} + \frac{1}{2} \left[1 - \frac{\sigma_1^2}{\sigma_0^2} \right] \exp\left(-\left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (70)$$

$$= \frac{1}{2} - \frac{1}{2} \left(\frac{\sigma_1^2}{\sigma_0^2} - 1 \right) \exp\left(-\left(1 - \frac{\sigma_0^2}{\sigma_1^2}\right)^{-1} \ln \frac{\sigma_1^2}{\sigma_0^2}\right) \quad (71)$$

$$(72)$$

(c)

Compare the optimal classifier to the QDA classifier in Braga-Neto (2020, Example 4.3). Compute the error of the QDA classifier and compare to the Bayes error. (Given $\sigma_0^2 = 2$ and $\sigma_1^2 = 8$)⁵

Part I: Optimal Error of Example 4.3

```

1 def berror_two(sig0, sig1):
2     assert sig1 > sig0
3     rat = sig1/sig0
4     return 0.5 - 0.5*(rat-1)*np.exp(-((1-rat**-1)**-1)*np.log(rat))
5
6
7 pd.DataFrame({"Optimal Error": [berror_two(2, 8)]})

```

5

For Problem 4.3(c), please assume $\sigma_0^2 = 2$ and $\sigma_1^2 = 8$. — [Ulisses \(TAMU Slack\)](#)

Optimal Error	
0	0.263765

Part II: QDA Error

Use the result in [Problem 4.4 \(b\)](#) and let \hat{r} be the boundary of the QDA in Braga-Neto (2020, Example 4.3):

- ϵ^0 is [6](#)

$$\epsilon^0 = \int_{\theta=0}^{\theta=2\pi} \int_{r=\hat{r}}^{\infty} \frac{1}{2\pi\hat{\sigma}_0^2} \exp\left(-\frac{r^2}{2\hat{\sigma}_0^2}\right) r dr d\theta \quad (73)$$

$$= \exp\left(-\frac{\hat{r}^2}{2\hat{\sigma}_0^2}\right) \quad (74)$$

$$= \exp\left(-\frac{\frac{32}{9} \ln 2}{2 \cdot \frac{2}{3}}\right) \quad (75)$$

$$\approx 0.157 \quad (76)$$

- ϵ^1 is [7](#)

$$\epsilon^1 = 1 - \exp\left(-\frac{\hat{r}^2}{2\hat{\sigma}_1^2}\right) \quad (77)$$

$$= 1 - \exp\left(-\frac{\frac{32}{9} \ln 2}{2 \cdot \frac{8}{3}}\right) \quad (78)$$

$$\approx 0.370 \quad (79)$$

Since $k_n = 0$ is assumed, the error of LDA is^{[8](#)}

$$\epsilon_{LDA} = \frac{1}{2}(\epsilon_{LDA}^0 + \epsilon_{LDA}^1) \quad (80)$$

$$= \frac{1}{2}(0.157 + 0.370) \quad (81)$$

$$= \underline{0.264} \quad (82)$$

Conclusion

The QDA error is larger than the optimal error.

⁶Via [WolframAlpha](#)

⁷Via [WolframAlpha](#)

⁸Via [WolframAlpha](#)

Problem 4.8 (Python Assignment)

Apply linear discriminant analysis to the stacking fault energy (SFE) dataset (see Braga-Neto (2020, sec. A8.4)), already mentioned in Braga-Neto (2020, ch. 1). Categorize the SFE values into two classes, low ($\text{SFE} \leq 35$) and high ($\text{SFE} \geq 45$), excluding the middle values.

(a)

Apply the preprocessing steps in `c01_matex.py` to obtain a data matrix of dimensions $123(\text{number of sample points}) \times 7(\text{number of features})$, as described in Braga-Neto (2020, sec. 1.8.2). Define low ($\text{SFE} \leq 35$) and high ($\text{SFE} \geq 45$) labels for the data. Pick the first 50% of the sample points to be the training data and the remaining 50% to be test data⁹.

```
1 # Setting
2 def get_SFE_low(df):
3     df_ = df[df["SFE"]<=35]
4     return df_.drop("SFE", axis=1)
5 def get_SFE_high(df):
6     df_ = df[df["SFE"]>=45]
7     return df_.drop("SFE", axis=1)
8 # Load data
9 SFE_data = pd.read_table("data/Stacking_Fault_Energy_Dataset.txt")
10 # pre-process the data
11 f_org = SFE_data.columns[:-1] # original features
12 n_org = SFE_data.shape[0] # original number of training points
13 p_org = np.sum(SFE_data.iloc[:, :-1]>0)/n_org # fraction of nonzero components for each fe
14 f_drp = f_org[p_org<0.6] # features with less than 60% nonzero co
15 SFE1 = SFE_data.drop(f_drp,axis=1) # drop those features
16 s_min = SFE1.min(axis=1)
17 SFE2 = SFE1[s_min!=0] # drop sample points with any zero values
18 SFE = SFE2[(SFE2.SFE<35)|(SFE2.SFE>45)] # drop sample points with middle responses
19 Y = SFE > 40
20 train, test, ytrain, ytest = train_test_split(SFE, Y, test_size=0.5, shuffle=False)

1 print(SFE.shape)
2 SFE.head(4)
```

(123, 8)

9

All, for the last problem, please make sure you are dividing the data 50% - 50% for training and testing, the values in the book are incorrect. — [Ulisses \(TAMU Slack\)](#)

Table 2: Filtered data

	C	N	Ni	Fe	Mn	Si	Cr	SFE
0	0.004	0.003	15.6	64.317	0.03	0.02	17.5	51.6
1	0.020	0.009	15.6	64.188	0.03	0.03	17.6	54.6
2	0.020	0.002	14.0	66.409	0.03	0.01	17.1	50.3
3	0.005	0.001	15.6	63.866	0.19	0.01	17.7	52.8

Table 3: Train data

	C	N	Ni	Fe	Mn	Si	Cr	SFE
0	0.004	0.003	15.6	64.317	0.03	0.02	17.5	51.6
1	0.020	0.009	15.6	64.188	0.03	0.03	17.6	54.6
2	0.020	0.002	14.0	66.409	0.03	0.01	17.1	50.3
3	0.005	0.001	15.6	63.866	0.19	0.01	17.7	52.8

```

1 print(train.shape)
2 train.head(4)

```

(61, 8)

```

1 print(test.shape)
2 test.head(4)

```

(62, 8)

(b)

Using the function `ttest_ind` from the `scipy.stats` module, apply Welch's two-sample t-test on the training data, and produce a table with the predictors, T statistic, and p -value, ordered with **largest absolute T** statistics at the top.

- For calculate independent t-test, we need to use `equal_var=False` for two

Table 4: Test data

	C	N	Ni	Fe	Mn	Si	Cr	SFE
295	0.07	0.40	16.13	54.818	9.64	0.45	18.48	65.0
296	0.07	0.54	16.13	54.678	9.64	0.45	18.48	53.0
297	0.04	0.04	9.00	70.920	1.20	0.40	18.20	30.4
298	0.04	0.04	9.00	70.920	1.20	0.40	18.20	25.7

Table 5: Results of T-test analysis

	Features	T-statistics	P-Values
2	Ni	-10.020338	9.325882e-14
3	Fe	6.050509	1.327194e-07
0	C	1.912194	6.242173e-02
5	Si	1.092405	2.816877e-01
1	N	-0.889981	3.789871e-01
4	Mn	-0.198507	8.436217e-01
6	Cr	0.038242	9.696653e-01

different sample sizes¹⁰.

```

1 df_low = get_SFE_low(train)
2 df_high = get_SFE_high(train)
3
4 ttest = st.ttest_ind(df_low, df_high, equal_var=False)
5
6 tdf = pd.DataFrame({
7     "Features": df_low.keys(),
8     "T-statistics": ttest.statistic,
9     "P-Values": ttest.pvalue
10 })
11
12 tdf = tdf.sort_values(["T-statistics"], ascending=[0], key=abs)
13 tdf

```

(c)

Pick the top two predictors and design an LDA classifier. (This is an example of *filter feature selection*, to be discussed in Chapter 9.). Plot the training data with the superimposed LDA decision boundary. Plot the testing data with the superimposed previously-obtained LDA decision boundary. Estimate the classification error rate on the training and test data. What do you observe?

Both train and test data has higher error rate compared to the result in Table 6. Because the train data (Figure 3, Figure 4) is inseparable with single boundary. This implies that we need extra informative feature to make this two classes separatable.

¹⁰https://docs.scipy.org/doc/scipy/reference/generated/scipy.stats.ttest_ind.html#r3566833beaa2-

```

1 features = list(tdf["Features"][0:2])
2 var1, var2 = features

1 def get_data_with_features(data, features):
2     X = data[features].values
3     Y = data.SFE > 40
4     return X, Y
5
6 def get_loss(X, Y, clf):
7     loss = sk.metrics.zero_one_loss(Y, clf.predict(X))
8     return loss
9
10 X, Y = get_data_with_features(train, features)
11 X_test, Y_test = get_data_with_features(test, features)
12
13 clf = LDA(priors=(0.5,0.5))
14 clf.fit(X,Y.astype(int))
15 a = clf.coef_[0]
16 b = clf.intercept_[0];
17
18 # Error
19 err_train = get_loss(X,Y, clf)
20 err_test = get_loss(X_test, Y_test, clf)

1 def plot_swe_predict(ax, X, Y, clf, title=""):
2     ax.scatter(X[~Y,0],X[~Y,1],c='blue',s=32,label='Low SFE')
3     ax.scatter(X[Y,0],X[Y,1],c='orange',s=32,label='High SFE')
4     left,right = ax.get_xlim()
5     bottom,top = ax.get_ylim()
6     ax.plot([left,right],[-left*a[0]/a[1]-b/a[1],-right*a[0]/a[1]-b/a[1]],'k',linewidth=2)
7     ax.set_title(title)
8     ax.set_xlim(left,right)
9     ax.set_ylim(bottom,top)
10    ax.set_xlabel(var1)
11    ax.set_ylabel(var2)
12    ax.legend();
13
14 fig43tr, ax43tr = plt.subplots()
15 ax43tr = plot_swe_predict(ax43tr, X, Y, clf, title="Train error {}".format(err_train));

```

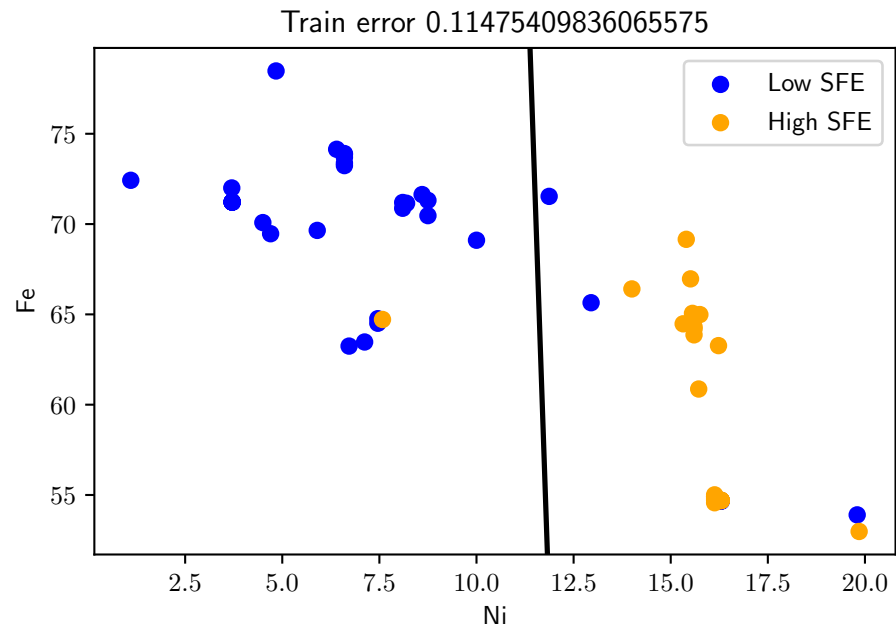


Figure 3: Train data with LDA

```

1 fig43t, ax43t = plt.subplots()
2 ax43t = plot_swe_predict(ax43t, X_test, Y_test, clf, title="Test error {}".format(err_test))

```




Figure 4: Test data with LDA

(d)

Repeat for the top three, and five predictors. Estimate the errors on the training and testing data (there is no need to plot the classifiers). What can you observe?

As shown in Table 6, the error rate does not decrease as more features get involved. This is because the sign of t-statistics are mixing together, and deviate the boundary to the optimal solution (see Table 5). The LDA approach may not appropriate for features with inversing T-statistics.

```

1  n_features = [2,3,4,5]
2  err_trains = np.zeros(len(n_features))
3  err_tests = np.zeros(len(n_features))
4  for (j,i) in enumerate(n_features):
5      fs = list(tdf["Features"][0:i])
6      X, Y = get_data_with_features(train, fs)
7      X_test, Y_test = get_data_with_features(test, fs)
8
9      clf = LDA(priors=(0.5,0.5))
10     clf.fit(X,Y.astype(int))
11

```

Table 6: Surveying different number of features

	Number of Features	Training Error Rate	Testing Error Rate
0	2	0.114754	0.145161
1	3	0.081967	0.145161
2	4	0.081967	0.225806
3	5	0.081967	0.193548

```

12     # Error
13     err_trains[j] = get_loss(X,Y, clf)
14     err_tests[j] = get_loss(X_test, Y_test, clf)
15
16 pd.DataFrame({
17     "Number of Features": n_features,
18     "Training Error Rate": err_trains,
19     "Testing Error Rate": err_tests
20 })

```

References

- ardianumam. 2017. “Understanding Multivariate Gaussian, Gaussian Properties and Gaussian Mixture Model.” *Ardian Umam Blog*.
- Braga-Neto, Ulisses. 2020. *Fundamentals of Pattern Recognition and Machine Learning*. Springer.